

Custom Theorems and Lists in LaTeX

Part I

SDF

Chapter 1

Introduction

Go to this hidden section.

Theorem 1.3: Pythagoras Theorem

In a right-angled triangle, the square of the hypotenuse is equal to the sum of the squares of the other two sides.

$$a^2 = b^2 + c^2$$

Concept 2.1: Probability

The probability of an event is a measure of the likelihood that the event will occur.

Chapter 2

Images and Lists

2.1 Lists in LaTeX

2.1.1 Unordered List

- First item
- Second item
- Third item

2.1.2 Ordered List

1. First item
2. Second item
3. Third item

Chapter 3

Code Blocks

3.1 Simple Code (Verbatim)

```
printf("Hello World");
```

Chapter 4

Optimization Problem

4.1 Definition of Optimization Problem

Definition 1.1: Optimization Problem

In an **optimization problem**, we minimize or maximize a function value, possibly subject to constraints.

$$\begin{aligned} & \underset{\theta \in \mathbb{R}^p}{\text{minimize}} && f(\theta) \\ & \text{subject to} && h_1(\theta) = 0, \quad h_2(\theta) = 0, \quad \dots, \quad h_m(\theta) = 0, \\ & && g_1(\theta) \leq 0, \quad g_2(\theta) \leq 0, \quad \dots, \quad g_n(\theta) \leq 0. \end{aligned}$$

- Decision variable: θ
- Objective function: f
- Equality constraint: $h_i(\theta) = 0$ for $i = 1, \dots, m$
- Inequality constraint: $g_j(\theta) \leq 0$ for $j = 1, \dots, n$

In machine learning (ML), we often minimize a "loss", but sometimes we maximize the "likelihood". In any case, minimization and maximization are equivalent since

$$\text{maximize } f(\theta) \quad \Leftrightarrow \quad \text{minimize } -f(\theta).$$

Definition 1.2: Feasible Point and Constraints

$\theta \in \mathbb{R}^p$ is a **feasible point** if it satisfies all constraints:

$$\begin{array}{ll} h_1(\theta) = 0 & g_1(\theta) \leq 0 \\ \vdots & \vdots \\ h_m(\theta) = 0 & g_n(\theta) \leq 0 \end{array}$$

Optimization problem is **infeasible** if there is no feasible point.

An optimization problem with no constraint is called an **unconstrained optimization problem**. Optimization problems with constraints is called a **constrained optimization problem**.

Definition 1.3: Optimal Value and Solution

Optimal value of an optimization problem is

$$p^* = \inf \{ f(\theta) \mid \theta \in \mathbb{R}^n, \theta \text{ feasible} \}$$

- $p^* = \infty$ if problem is infeasible
- $p^* = -\infty$ is possible
- In ML, it is often a priori clear that $0 \leq p^* < \infty$.

If $f(\theta^*) = p^*$, we say θ^* is a **solution** or θ^* is **optimal**. A solution may or may not exist, and a solution may or may not be unique.

4.2 Examples of Optimization Problem

Example 1.4: Curve Fitting

Consider setup with data X_1, \dots, X_N and corresponding labels Y_1, \dots, Y_N satisfying the relationship

$$Y_i = f_*(X_i) + \text{error}$$

for $i = 1, \dots, N$. Hopefully, "error" is small. True function f_* is unknown. Goal is to find a function (curve) f such that $f \approx f_*$.

Example 1.5: Least-Squares Minimization

- **Problem**

$$\underset{\theta \in \mathbb{R}^p}{\text{minimize}} \quad \frac{1}{2} \|X\theta - Y\|^2$$

where $X \in \mathbb{R}^{N \times p}$ and $Y \in \mathbb{R}^N$. Equivalent to

$$\underset{\theta \in \mathbb{R}^p}{\text{minimize}} \quad \frac{1}{2} \sum_{i=1}^N (X_i^\top \theta - Y_i)^2$$

$$\text{where } X = \begin{bmatrix} X_1^\top \\ \vdots \\ X_N^\top \end{bmatrix} \text{ and } Y = \begin{bmatrix} Y_1 \\ \vdots \\ Y_N \end{bmatrix}.$$

-
- **Solution**

To solve

$$\underset{\theta \in \mathbb{R}^p}{\text{minimize}} \frac{1}{2} \|X\theta - Y\|^2$$

take gradient and set it to 0.

$$\nabla_{\theta} \frac{1}{2} \|X\theta - Y\|^2 = X^{\top} (X\theta - Y)$$

$$X^{\top} (X\theta^* - Y) = 0$$

$$\theta^* = (X^{\top} X)^{-1} X^{\top} Y$$

Here, we assume $X^{\top} X$ is invertible.

Concept 1.6: Least squares is an instance of curve fitting.

Define $f_{\theta}(x) = x^{\top} \theta$. Then LS becomes

$$\underset{\theta \in \mathbb{R}^p}{\text{minimize}} \frac{1}{2} \sum_{i=1}^N (f_{\theta}(X_i) - Y_i)^2$$

and the solution hopefully satisfies

$$Y_i = f_{\theta}(X_i) + \text{small.}$$

Since X_i and Y_i is assumed to satisfy

$$Y_i = f_{\star}(X_i) + \text{error}$$


we are searching over linear functions (linear curves) f_{θ} that best fit (approximate) f_{\star} .

4.3 Local and Global Minimum

Definition 1.7: Local vs Global Minima

θ^* is a ****local minimum**** if $f(\theta) \geq f(\theta^*)$ for all feasible θ within a small neighborhood.

θ^* is a ****global minimum**** if $f(\theta) \geq f(\theta^*)$ for all feasible θ .



In the worst case, finding the global minimum of an optimization problem is difficult. However, in deep learning, optimization problems are often "solved" without any guarantee of global optimality.

Chapter 5

Basics of Monte Carlo

5.1 Monte Carlo Estimation

Definition 13.1: Monte Carlo Estimation

Consider IID data $X_1, \dots, X_N \sim f$. Let $\phi(X) \geq 0$ be some function. (The assumption $\phi(X) \geq 0$ can be relaxed.) Consider the problem of estimating

$$I = \mathbb{E}_{X \sim f}[\phi(X)] = \int \phi(x)f(x)dx$$

One commonly uses

$$\hat{I}_N = \frac{1}{N} \sum_{i=1}^N \phi(X_i)$$

to estimate I , which is called **monte carlo estimation**. After all, $\mathbb{E}[\hat{I}_N] = I$ and $\hat{I}_N \rightarrow I$ by the law of large numbers. (Convergence in probability by weak law of large numbers and almost sure convergence by strong law of large numbers.)

Concept 13.2: Evidence of Convergence for Monte Carlo Estimation

We can quantify convergence with variance:

$$\text{Var}_{X \sim f}(\hat{I}_N) = \sum_{i=1}^N \text{Var}_{X_i \sim f}\left(\frac{\phi(X_i)}{N}\right) = \frac{1}{N} \text{Var}_{X \sim f}(\phi(X))$$

In other words

$$\mathbb{E}\left[\left(\hat{I}_N - I\right)^2\right] = \frac{1}{N} \text{Var}_{X \sim f}(\phi(X))$$

and

$$\mathbb{E} \left[\left(\hat{I}_N - I \right)^2 \right] \rightarrow 0$$

as $N \rightarrow \infty$. So, $\hat{I}_N \rightarrow I$ in L^2 provided that $\text{Var}_{X \sim f}(\phi(X)) < \infty$.

Definition 13.3: Empirical Risk Minimization (ERM)

In machine learning and statistics, we often wish to solve

$$\underset{\theta \in \Theta}{\text{minimize}} \quad \mathcal{L}(\theta)$$

where the objective function

$$\mathcal{L}(\theta) = \mathbb{E}_{X \sim p_X} [\ell(f_\theta(X), f_\star(X))]$$

Is the (true) risk. However, the evaluation of $\mathbb{E}_{X \sim p_X}$ is impossible (if p_X is unknown) or intractable (if p_X is known but the expectation has no closed-form solution). Therefore, we define the proxy loss function

$$\mathcal{L}_N(\theta) = \frac{1}{N} \sum_{i=1}^N \ell(f_\theta(X_i), f_\star(X_i))$$

which we call the empirical risk, and solve

$$\underset{\theta \in \Theta}{\text{minimize}} \quad \mathcal{L}_N(\theta)$$

This is called **empirical risk minimization (ERM)**. The idea is that

$$\mathcal{L}_N(\theta) \approx \mathcal{L}(\theta)$$

with high probability, so minimizing $\mathcal{L}_N(\theta)$ should be similar to minimizing $\mathcal{L}(\theta)$.

Concept 13.4: Evidence of Convergence for Empirical Risk Minimization

Technical note) The law of large numbers tells us that

$$\mathbb{P}(|\mathcal{L}_N(\theta) - \mathcal{L}(\theta)| > \varepsilon) = \text{small}$$

for any given θ , but we need

$$\mathbb{P} \left(\sup_{\theta \in \Theta} |\mathcal{L}_N(\theta) - \mathcal{L}(\theta)| > \varepsilon \right) = \text{small}$$

for all compact Θ in order to conclude that the argmins of the two losses to be similar. These types of results are established by a uniform law of large numbers.

5.2 Importance Sampling

Definition 13.5: Importance Sampling (IS)

Importance sampling (IS) is a technique for reducing the variance of a Monte Carlo estimator.

Key insight of important sampling:

$$I = \mathbb{E}_{X \sim f}[\phi(X)] = \int \phi(x)f(x)dx = \int \frac{\phi(x)f(x)}{g(x)}g(x)dx = \mathbb{E}_{X \sim g}\left[\frac{\phi(X)f(X)}{g(X)}\right]$$

(We do have to be mindful of division by 0.) Then

$$\hat{I}_N = \frac{1}{N} \sum_{i=1}^N \phi(X_i) \frac{f(X_i)}{g(X_i)}$$

with $X_1, \dots, X_N \sim g$ is also an estimator of I . Indeed, $\mathbb{E}[\hat{I}_N] = I$ and $\hat{I}_N \rightarrow I$. The weight $\frac{f(x)}{g(x)}$ is called the **likelihood ratio** or the **Radon-Nikodym derivative**.

So we can use samples from g to compute expectation with respect to f .

Example 13.6: IS Example

Consider the setup of estimating the probability

$$\mathbb{P}(X > 3) = 0.00135$$

where $X \sim \mathcal{N}(0, 1)$. If we use the regular Monte Carlo estimator

$$\hat{I}_N = \frac{1}{N} \sum_{i=1}^N \mathbf{1}_{\{X_i > 3\}}$$

where $X_i \sim \mathcal{N}(0, 1)$, if N is not sufficiently large, we can have $\hat{I}_N = 0$. Inaccurate estimate.

If we use the IS estimator

$$\hat{I}_N = \frac{1}{N} \sum_{i=1}^N \mathbf{1}_{\{Y_i > 3\}} \exp\left(\frac{(Y_i - 3)^2 - Y_i^2}{2}\right)$$

where $Y_i \sim \mathcal{N}(3, 1)$, having $\hat{I}_N = 0$ is much less likely. Estimate is much more accurate.

Concept 13.7: Optimal Sampling Distribution

Benefit of IS quantified by with variance:

$$\text{Var}_{X \sim g}(\hat{I}_N) = \sum_{i=1}^N \text{Var}_{X \sim g} \left(\frac{\phi(X_i) f(X_i)}{ng(X_i)} \right) = \frac{1}{N} \text{Var}_{X \sim g} \left(\frac{\phi(X) f(X)}{g(X)} \right)$$

If $\text{Var}_{X \sim g} \left(\frac{\phi(X) f(X)}{g(X)} \right) < \text{Var}_{X \sim f}(\phi(X))$, then IS provides variance reduction. We call g the importance or sampling distribution. Choosing g poorly can increase the variance. What is the best choice of g ?

The sampling distribution

$$g(x) = \frac{\phi(x)f(x)}{I}$$

makes $\text{Var}_{X \sim g} \left(\frac{\phi(X)f(X)}{g(X)} \right) = \text{Var}_{X \sim g}(I) = 0$ and therefore is optimal. (I serves as the normalizing factor that ensures the density g integrates to 1.) Problem: Since we do not know the normalizing factor I , the answer we wish to estimate, sampling from g is usually difficult.

Concept 13.8: Optimized / Trained Sampling Distribution

Instead, we consider the optimization problem

$$\underset{g \in \mathcal{G}}{\text{minimize}} \quad D_{\text{KL}} \left(g \parallel \frac{\phi f}{I} \right)$$

and compute a suboptimal, but good, sampling distribution within a class of sampling distributions \mathcal{G} . (In ML, $\mathcal{G} = \{g_\theta \mid \theta \in \Theta\}$ is parameterized by neural networks.)

Importantly, this optimization problem does not require knowledge of I .

$$\begin{aligned} D_{\text{KL}}(g_\theta \parallel \phi f / I) &= \mathbb{E}_{X \sim g_\theta} \left[\log \left(\frac{I g_\theta(X)}{\phi(X) f(X)} \right) \right] \\ &= \mathbb{E}_{X \sim g_\theta} \left[\log \left(\frac{g_\theta(X)}{\phi(X) f(X)} \right) \right] + \log I \\ &= \mathbb{E}_{X \sim g_\theta} \left[\log \left(\frac{g_\theta(X)}{\phi(X) f(X)} \right) \right] + \text{constant independent of } \theta \end{aligned}$$

How do we compute stochastic gradients?

5.3 Log-Derivative Trick

Definition 13.9: Log-Derivative Trick

Generally, consider the setup where we wish to solve

$$\underset{\theta \in \mathbb{R}^p}{\text{minimize}} \mathbb{E}_{X \sim f_\theta} [\phi(X)]$$

with SGD. (Previous situation (Concept 13.8) had θ -dependence both on and inside the expectation. For now, let's simplify the problem so that ϕ does not depend on θ .)

Incorrect gradient computation:

$$\nabla_\theta \mathbb{E}_{X \sim f_\theta} [\phi(X)] \stackrel{?}{=} \mathbb{E}_{X \sim f_\theta} [\nabla_\theta \phi(X)] = \mathbb{E}_{X \sim f_\theta} [0] = 0$$

Correct gradient computation:

$$\begin{aligned} \nabla_\theta \mathbb{E}_{X \sim f_\theta} [\phi(X)] &= \nabla_\theta \int \phi(x) f_\theta(x) dx = \int \phi(x) \nabla_\theta f_\theta(x) dx \\ &= \int \phi(x) \frac{\nabla_\theta f_\theta(x)}{f_\theta(x)} f_\theta(x) dx = \mathbb{E}_{X \sim f_\theta} \left[\phi(X) \frac{\nabla_\theta f_\theta(X)}{f_\theta(X)} \right] \\ &= \mathbb{E}_{X \sim f_\theta} [\phi(X) \nabla_\theta \log(f_\theta(X))] \end{aligned}$$

Therefore, $\phi(X) \nabla_\theta \log(f_\theta(X))$ with $X \sim f_\theta$ is a stochastic gradient of the loss function. This technique is called the log-derivative trick, the likelihood ratio gradient[#], or REINFORCE^{*}.

Formula with the log-derivative ($\nabla_\theta \log(\cdot)$) is convenient when dealing with Gaussians, or more generally exponential families, since the densities are of the form

$$f_\theta(x) = h(x) \exp(\text{function of } \theta)$$

([#]P. W. Glynn, Likelihood ratio gradient estimation for stochastic systems, Communications of the ACM, 1990.

^{*}R. J. Williams, Simple statistical gradient-following algorithms for connectionist reinforcement learning. Machine Learning, 1992.)

Example 13.10: Log-Derivative Trick Example

Learn $\mu \in \mathbb{R}^2$ to minimize the objective below.

$$\underset{\mu \in \mathbb{R}^2}{\text{minimize}} \mathbb{E}_{X \sim \mathcal{N}(\mu, I)} \left\| X - \begin{pmatrix} 5 \\ 5 \end{pmatrix} \right\|^2$$


Then the loss function is

$$\mathcal{L}(\mu) = \mathbb{E}_{X \sim \mathcal{N}(\mu, I)} \left\| X - \begin{pmatrix} 5 \\ 5 \end{pmatrix} \right\|^2 = \int \left\| x - \begin{pmatrix} 5 \\ 5 \end{pmatrix} \right\|^2 \frac{1}{2\pi} \exp\left(-\frac{1}{2}\|x - \mu\|^2\right) dx$$

And, using $X_1, \dots, X_B \sim \mathcal{N}(\mu, I)$, we have stochastic gradients

$$\nabla_{\mu} \mathcal{L}(\mu) = \mathbb{E}_{X \sim \mathcal{N}(\mu, I)} \left[\left\| x - \begin{pmatrix} 5 \\ 5 \end{pmatrix} \right\|^2 \nabla_{\mu} \left(-\frac{1}{2} \|x - \mu\|^2 \right) \right] \approx \frac{1}{B} \sum_{i=1}^B \left\| X_i - \begin{pmatrix} 5 \\ 5 \end{pmatrix} \right\|^2 (X_i - \mu)$$

These stochastic gradients have large variance and thus SGD is slow.



5.4 Reparameterization Trick

Definition 13.11: Reparameterization Trick

The reparameterization trick (RT) or the pathwise derivative (PD) relies on the key insight.

$$\mathbb{E}_{X \sim \mathcal{N}(\mu, \sigma^2)} [\phi(X)] = \mathbb{E}_{Y \sim \mathcal{N}(0,1)} [\phi(\mu + \sigma Y)]$$

Gradient computation:

$$\begin{aligned} \nabla_{\mu, \sigma} \mathbb{E}_{X \sim \mathcal{N}(\mu, \sigma^2)} [\phi(X)] &= \mathbb{E}_{Y \sim \mathcal{N}(0,1)} [\nabla_{\mu, \sigma} \phi(\mu + \sigma Y)] = \mathbb{E}_{Y \sim \mathcal{N}(0,1)} \left[\phi'(\mu + \sigma Y) \begin{bmatrix} 1 \\ Y \end{bmatrix} \right] \\ &\approx \frac{1}{B} \sum_{i=1}^B \phi'(\mu + \sigma Y_i) \begin{bmatrix} 1 \\ Y_i \end{bmatrix}, \quad Y_1, \dots, Y_B \sim \mathcal{N}(0, I) \end{aligned}$$

RT is less general than log-derivative trick, but it usually produces stochastic gradients with lower variance.

Example 13.12: Reparameterization Trick Example

Consider the same example as before

$$\mathcal{L}(\mu) = \mathbb{E}_{X \sim \mathcal{N}(\mu, I)} \left\| X - \begin{pmatrix} 5 \\ 5 \end{pmatrix} \right\|^2 = \mathbb{E}_{Y \sim \mathcal{N}(0, I)} \left\| Y + \mu - \begin{pmatrix} 5 \\ 5 \end{pmatrix} \right\|^2$$

Gradient computation:

$$\begin{aligned}\nabla_{\mu}\mathcal{L}(\mu) &= \mathbb{E}_{Y \sim \mathcal{N}(0, I)} \nabla_{\mu} \left\| Y + \mu - \begin{pmatrix} 5 \\ 5 \end{pmatrix} \right\|^2 = 2\mathbb{E}_{Y \sim \mathcal{N}(0, I)} \left(Y + \mu - \begin{pmatrix} 5 \\ 5 \end{pmatrix} \right) \\ &\approx \frac{2}{B} \sum_{i=1}^B \left(Y_i + \mu - \begin{pmatrix} 5 \\ 5 \end{pmatrix} \right), \quad Y_1, \dots, Y_B \sim \mathcal{N}(0, I)\end{aligned}$$

These stochastic gradients have smaller variance and thus SGD is faster.

Example 13.13: Log Derivative Trick vs Reparameterization Trick

The image below is the result of SGD with the computed gradients by Example 13.10 and Example 13.12.

